

# Exploring Distributed Computing

Distributed computing allows multiple computers to work together on a single task, enhancing performance and resilience. Let's delve into its key aspects.

## Components of Distributed Computing



### Nodes

Individual computers or processors that are part of the distributed system, performing computational tasks.



### Network

The communication infrastructure connecting the nodes, enabling data exchange and coordination between them.



### Middleware

Software layer that sits between the operating system and applications, managing communication and resource sharing.

## Benefits of Distributed Computing



### Scalability

Easily expand computational power by adding more nodes, adapting to increasing workload demands.



### Reliability

Reduces the risk of single points of failure, as tasks can be redistributed if one node goes down.



### Efficiency

Processes tasks faster by parallelizing them across multiple machines, optimizing resource utilization.

# Microservices: The Real-World Scenario

Imagine you are building a machine learning system to recommend movies to users (like Netflix). This ML system has several critical components:



## Data Ingestion

Collects and processes data from users (e.g., their watch history and ratings).



## Feature Engineering

Transforms raw data into meaningful inputs for your ML model.



## Model Training

Continuously trains and updates your recommendation algorithm.



## Model Serving

Hosts the trained model and responds to user requests in real time.



## User Interface

Provides the website or app where users can browse and see recommendations.

In the traditional **monolithic architecture**, all these components would be bundled into a single application. This means if you wanted to scale the system (e.g., due to an increase in user requests), you would have to scale the entire application.

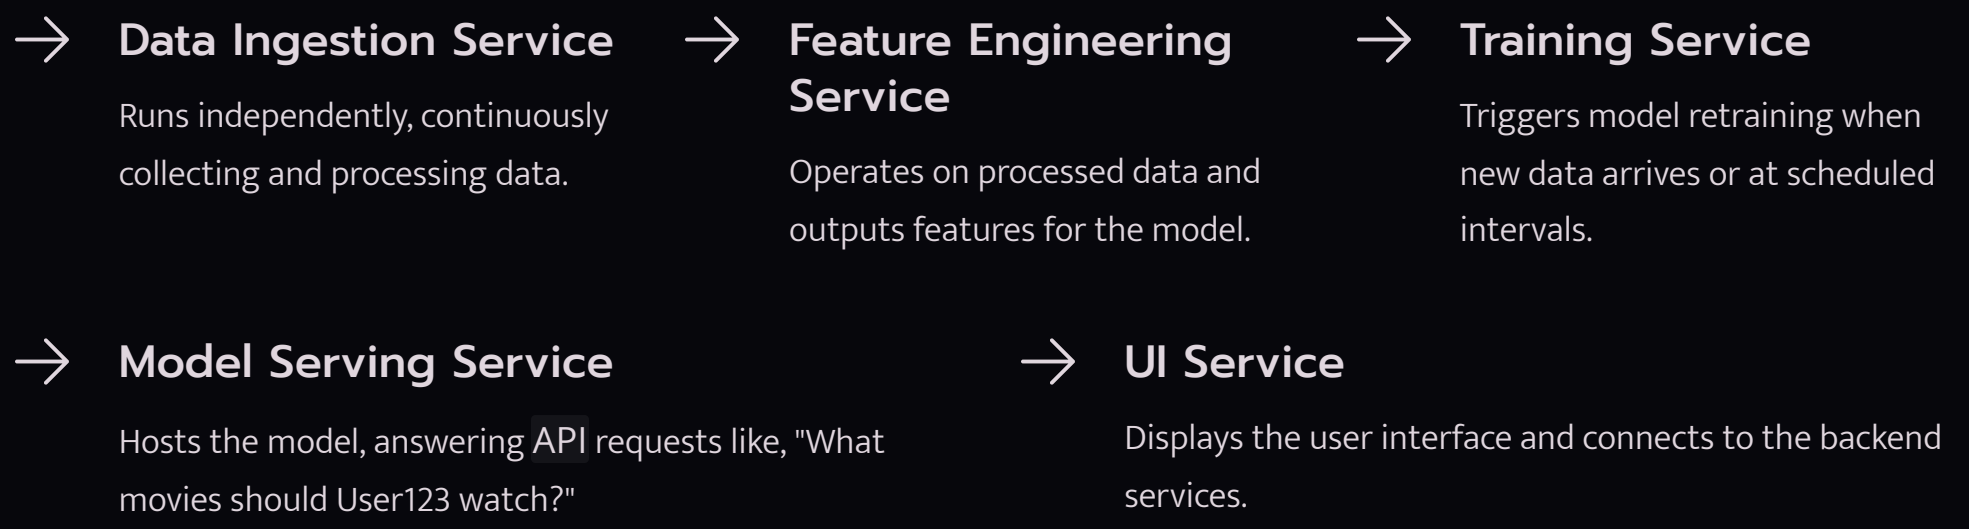
This can be inefficient, especially if only a specific component, like **Model Serving**, needs more resources.

Now, let's see how **microservices** solve this problem.

# Microservices in Action

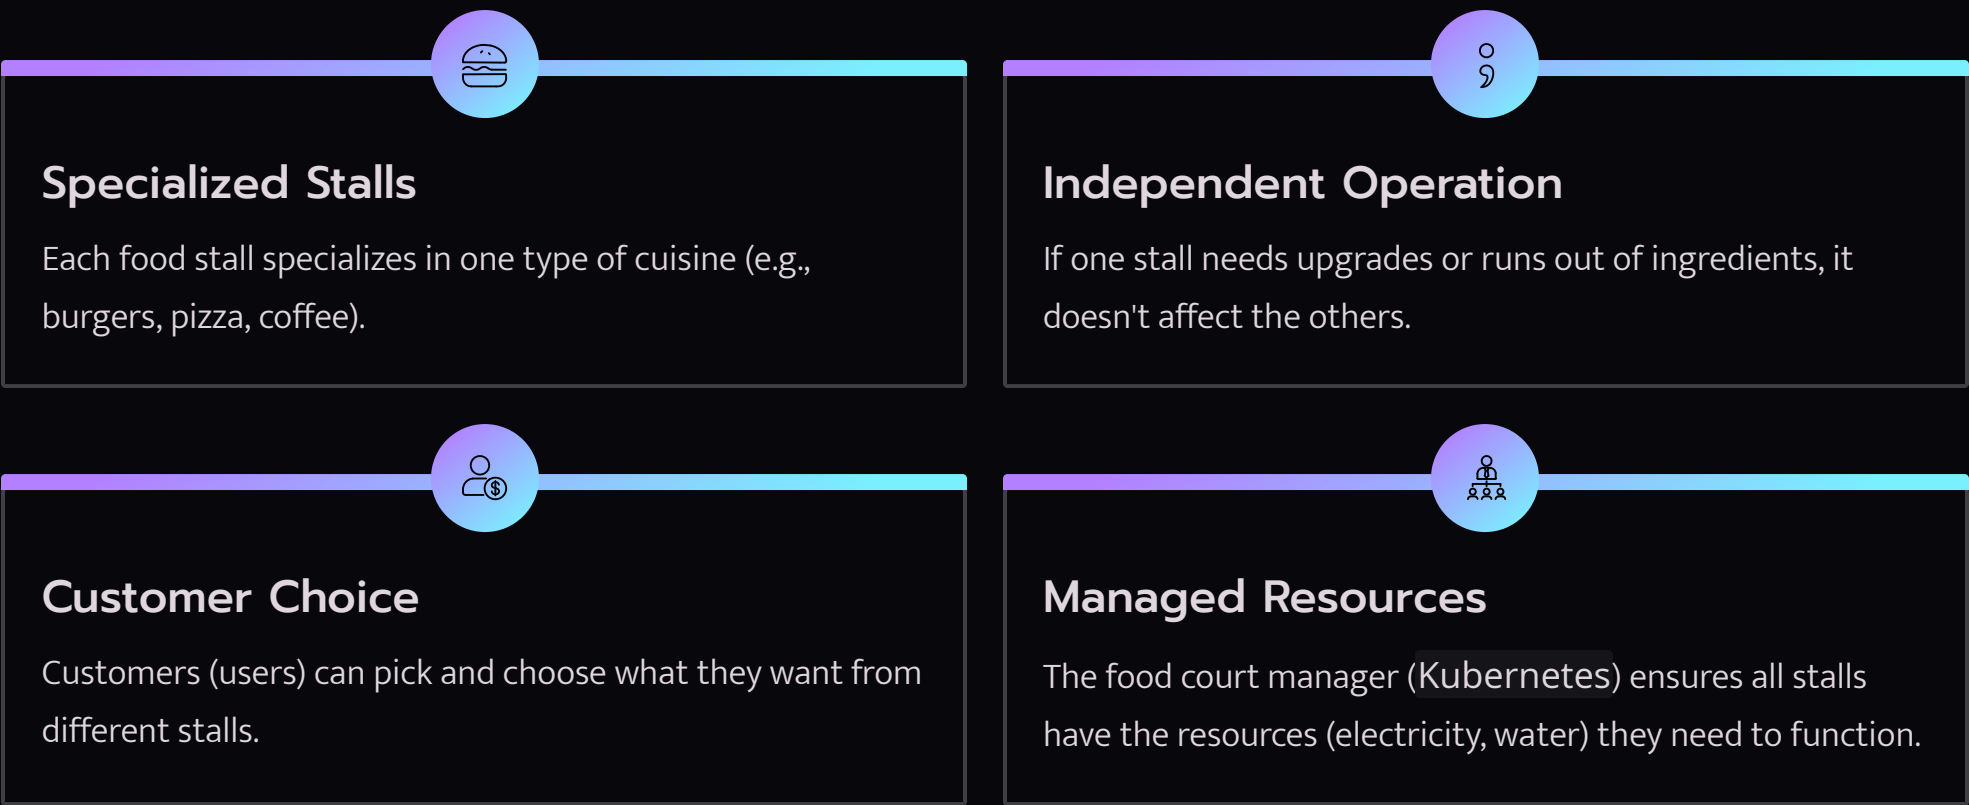
Instead of bundling everything together, **microservices** break down an application into smaller, independent components. Each component is responsible for a single task and can run, scale, and be updated independently.

For our ML system, this means separating the core functions into distinct services:



This separation makes it easy to scale just the **Model Serving Service** when the number of user requests spikes, without affecting the other parts of the system.

## Real-Life Analogy: The Food Court



## Challenges of Distributed Computing

While microservices offer flexibility, managing many independent services introduces complexities. This is where **Kubernetes** becomes essential.

### How Kubernetes Addresses These Challenges

**Kubernetes** is a container orchestration platform specifically designed to address the challenges of distributed computing. It automates much of the manual effort involved in deploying, managing, and scaling containerized applications.



# **\*\* Kubernetes Internals \*\***

**Master Node (Control Plane):** The master node is the brain of the Kubernetes cluster. It oversees the system, manages workloads, and ensures everything runs as expected. Think of it as the manager in a factory that delegates tasks and monitors operations.

**Resource Manager:** The resource manager in Kubernetes ensures that cluster resources like CPU, memory, and storage are allocated efficiently. It's like a warehouse manager who ensures that raw materials (resources) are distributed to the right production lines (nodes).

**API Server (For user comm):** Users interact with Kubernetes through the API Server. It's like the receptionist at an office—you send requests (like deploying an app), and the API Server routes them to the appropriate components in the cluster.

**Database (etcd):** etcd is Kubernetes' central database where all cluster data is stored, including the current state of the system and desired configurations. Imagine it as a library catalog—every time you check out or return a book, the catalog is updated to reflect the changes.

**Worker Node:** Worker nodes are the muscle of the cluster. They run your applications and handle the tasks assigned by the master node. Think of them as factory workers who execute tasks given by the manager (master node).

**Kubelet:** The kubelet is an agent running on each worker node that ensures containers (your applications) are running as expected. It's like a shift supervisor in a factory who ensures each machine is operating correctly.

**Kube-Proxy:** The kube-proxy manages network traffic for pods, ensuring they can communicate with each other and the outside world. It's like a traffic cop at a busy intersection, directing cars (data packets) to the right destinations.

**Pods:** A pod is the smallest deployable unit in Kubernetes and typically wraps one or more containers. Think of it as a container ship holding one or more goods (containers) and transporting them across a logistics network.

**SharedDB (Volumes):** Volumes are shared storage spaces in Kubernetes that allow pods to save and share data. It's like a shared locker room where workers (pods) can access tools or leave notes for each other.

**Kube-Manifest (YAML):** Kube manifests are configuration files written in YAML that define what you want Kubernetes to do, such as deploying an app or creating a service. Think of it as the blueprint for building a house—Kubernetes reads it to know exactly what to construct.

**Service:** A service in Kubernetes provides a stable network endpoint for accessing a set of pods. Even if the pods are replaced or moved, the service ensures they can still be reached. It's like a restaurant hotline—no matter who answers the phone (which server pod), your order is taken.

**Namespace:** Namespaces are virtual clusters within a Kubernetes cluster that help organize and isolate resources. It's like different departments in a large office building—HR, Sales, and IT each work in separate areas but share the same infrastructure.

**Scheduler:** Decides which worker node will run a new pod based on resource availability. It's like a dispatcher allocating tasks to the most available worker.

**ReplicaSets:** Ensure that a specified number of identical pods are always running. It's like a backup generator ensuring there's always power even if one generator fails.

## **\*\*\*\*\* Real-Life Analogy: Distributed Computing Without and With Kubernetes \*\*\*\*\***

### **- Without Kubernetes:**

Imagine running a massive restaurant chain where you have to manage chefs, waiters, supplies, and customer orders manually for each branch. If one branch runs out of ingredients or if a waiter quits, everything falls apart.

### **- With Kubernetes:**

With Kubernetes, it's like having an automated restaurant management system. It automatically handles ingredients (resources), assigns tasks to chefs (pods), and ensures customer orders (requests) are fulfilled even if some chefs are busy or new ones are added. It provides the resilience and scalability needed for a complex distributed system.

## **\*\*Summary\*\***

Kubernetes simplifies the management of complex distributed systems by automating key operational tasks, ensuring high availability, and enabling efficient resource utilization. It transforms a chaotic manual process into an orchestrated, self-healing system.

### **For example:**

If a particular service (like the "Model Serving Service" from the previous card) experiences a spike in demand, Kubernetes can automatically scale up the number of pods running that service to handle the load, and then scale them back down when demand subsides. This dynamic scaling is crucial for maintaining performance and cost-efficiency in a microservices architecture.

## **Summary**

This analogy highlights how Kubernetes acts as the central orchestrator, turning a collection of independent, potentially fragile microservices into a robust, scalable, and manageable system, much like an efficient manager streamlines a complex operation.